
Data preparation

The code for this part is left out from the PDF due to its length...

Data Cleaning

SVM Cross Validation

```
set.seed(123)
cols <- c("Attrition", "EnvironmentSatisfaction", "NumCompaniesWorked", "JobSatisfaction", "BusinessTravel")

emp_data <- emp_attrition[cols]

#Split
split <- initial_split(emp_data, prop = 0.7, strata = 'Attrition')
#Create training and test set
training <- training(split)
testing <- testing(split)


trctrl <- trainControl(method = "repeatedcv", number=10, repeats = 3)
svm_radial <- train(Attrition ~ ., data = training, method = "svmRadial", trControl=trctrl, preProcess = NULL)
svm_linear <- train(Attrition ~ ., data = training, method = "svmLinear", trControl=trctrl, preProcess = NULL)

test_pred_radial <- predict(svm_radial, newdata = testing)
test_pred_linear <- predict(svm_linear, newdata = testing)

cm_radial <- confusionMatrix(table(test_pred_radial, testing$Attrition))
cm_linear <- confusionMatrix(table(test_pred_linear, testing$Attrition))

draw_confusion_matrix(cm_radial)
```

CONFUSION MATRIX




DETAILS

| | | | | |
|-----------------------------|-----------------------------|---------------------------|------------------------|--------------------|
| Sensitivity 0.989 | Specificity 0.153 | Precision 0.857 | Recall 0.989 | F1 0.918 |
| Accuracy 0.853 | | Kappa 0.208 | | |

```
draw_confusion_matrix(cm_linear)
```

CONFUSION MATRIX



DETAILS

| | | | | |
|--------------------------|-------------------------|---------------------------|--------------------|--------------------|
| Sensitivity 1 | Specificity 0 | Precision 0.837 | Recall 1 | F1 0.911 |
| Accuracy 0.837 | | Kappa 0 | | |

Above two model shows that the model is quite biased towards predicting “No” Attrition. This can be assumed due to the unbalanced dataset where about 80% of the dataset included employees that did not have Attrition and 20% who had Attrition.

Even though the accuracy of two model is relatively high scoring 85% and 83% this model is not appropriate to predict employees that will have Attrition due to low specificity.

```
emp_data %>% count(Attrition)
```

```
##   Attrition    n
## 1         No 1233
## 2         Yes  237
```

```
new_emp<- emp_data %>% group_by(Attrition) %>% slice_sample(n=200)
```

```
#Split
```

```
split2 <- initial_split(new_emp, prop = 0.7, strata = 'Attrition')
```

```
#Create training and test set
```

```
training2 <- training(split2)
```

```
testing2 <- testing(split2)
```

```
trctrl <- trainControl(method = "repeatedcv", number=10, repeats = 3)
```

```
svm_radial2 <- train(Attrition ~ ., data = training2, method = "svmRadial", trControl=trctrl, preProcess=
```

```
svm_linear2 <- train(Attrition ~ ., data = training2, method = "svmLinear", trControl=trctrl, preProcess=
```

```
test_pred_radial2 <- predict(svm_radial2, newdata = testing2)
```


```
test_pred_linear2 <- predict(svm_linear2, newdata = testing2)
```

```
cm_radial2 <- confusionMatrix(table(test_pred_radial2, testing2$Attrition))
```

```
cm_linear2 <- confusionMatrix(table(test_pred_linear2, testing2$Attrition))
```

```
draw_confusion_matrix(cm_radial2)
```

CONFUSION MATRIX




DETAILS

| | | | | |
|---------------------------|----------------------------|---------------------------|----------------------|--------------------|
| Sensitivity 0.7 | Specificity 0.75 | Precision 0.737 | Recall 0.7 | F1 0.718 |
| Accuracy 0.725 | | Kappa 0.45 | | |

```
draw_confusion_matrix(cm_linear2)
```

CONFUSION MATRIX



DETAILS

| | | | | |
|----------------------------|-----------------------------|--------------------------|-----------------------|--------------------|
| Sensitivity 0.65 | Specificity 0.817 | Precision 0.78 | Recall 0.65 | F1 0.709 |
| Accuracy 0.733 | | Kappa 0.467 | | |

After undersampling the dataset to have 50:50 ratio, the two model has lower accuracy but improved in specificity.


Model Tuning

Since we observed that dataset with 50:50 ratio showed desirable prediction output for our case, we will optimize the model by tuning the parameters.

```
set.seed(123)
grid <- expand.grid(C = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 5))

svm_linear_grid <- train(
  Attrition ~ ., data= training2, method = "svmLinear",
  trControl = trctrl,
  preProcess = c("center", "scale"),
  tuneGrid = grid,
  tuneLength = 10
)

plot(svm_linear_grid)
```



```
svm_linear_grid$bestTune
```

```
##      C
## 4 0.25
```


```

grid2 <- expand.grid(C = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 5), sigma=c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 5))

svm_radial_grid <- train(
  Attrition ~ ., data= training2, method = "svmRadial",
  trControl = trctrl,
  preProcess = c("center", "scale"),
  tuneGrid = grid2,
  tuneLength = 10
)

plot(svm_radial_grid)

```



```
svm_radial_grid$bestTune
```

```
##      sigma      C
## 97  0.01 1.5
```

The result shows C of 0.25 is the best parameter for linear SVM, for Radial SVM sigma of 0.01 and C of 1.5. Even though these values were obtained through 10 fold cross validation, the optimal parameter may change when the random seed is set to other value most likely because our dataset is not large enough and prone to random initialization of training and testing splitting.

Neural Network

```
library(neuralnet)
```

```
##
```

```
## Attaching package: 'neuralnet'
```

```
## The following object is masked from 'package:ROCR':
```

```
##
```

```
##      prediction
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      compute
```

```
set.seed(123)
```

```
emp_net = nnet(Attrition ~ ., data=training, size=15, maxit=100, range=0.1, decay=5e-4)
```

```
## # weights:  301
```

```
## initial  value 578.352568
```

```
## iter   10 value 422.029830
```

```
## iter   20 value 381.411616
```

```
## iter   30 value 344.878275
```

```
## iter   40 value 316.048645
```

```
## iter   50 value 290.011759
```

```
## iter   60 value 280.742910
```

```
## iter   70 value 277.412237
```

```
## iter   80 value 274.046773
```

```
## iter   90 value 269.007408
```

```
## iter  100 value 260.534699
```

```
## final   value 260.534699
```

```
## stopped after 100 iterations
```

```
pred_net <- predict(emp_net, testing, type="class")
```

```
cm_nn <- confusionMatrix(table(pred=pred_net, true=testing$Attrition))
```

```
draw_confusion_matrix(cm_nn)
```

CONFUSION MATRIX

| | | Actual | |
|-----------|-----|--------|-----|
| | | No | Yes |
| Predicted | No | 348 | 55 |
| | Yes | 22 | 17 |


DETAILS

| | | | | |
|-----------------------------|-----------------------------|---------------------------|------------------------|------------------|
| Sensitivity 0.941 | Specificity 0.236 | Precision 0.864 | Recall 0.941 | F1 0.9 |
| Accuracy 0.826 | | Kappa 0.217 | | |

```
pred_raw <- predict(emp_net, testing, decision.values=TRUE, type = "raw")
```

```
pred <- ROCR::prediction(pred_raw, testing$Attrition)
perf <- ROCR::performance(pred,measure="tpr", x.measure="fpr")
plot(perf, lwd=2, col="blue")
abline(a=0, b=1)
auc_ROCR <- performance(pred, measure="auc")
auc_ROCR <- auc_ROCR@y.values[[1]]
auc_ROCR
```

```
## [1] 0.7371246
```

```
emp_net2 = nnet(Attrition ~ ., data=training2, size=15, maxit=100, range=0.1, decay=5e-4)
```

```
## # weights: 301
## initial value 207.110397
## iter 10 value 181.489799
## iter 20 value 171.918737
## iter 30 value 155.710904
## iter 40 value 144.243328
## iter 50 value 134.161647
## iter 60 value 129.096882
## iter 70 value 128.286932
## iter 80 value 127.540918
## iter 90 value 127.163612
## iter 100 value 126.672924
## final value 126.672924
## stopped after 100 iterations
```

```
pred_net2 <- predict(emp_net2, testing2, type="class")
cm_nn2 <- confusionMatrix(table(pred=pred_net2, true=testing2$Attrition))
draw_confusion_matrix(cm_nn2)
```

CONFUSION MATRIX

| | | Actual | |
|-----------|-----|--------|-----|
| | | No | Yes |
| Predicted | No | 47 | 27 |
| | Yes | 13 | 33 |

DETAILS

| | | | | |
|-----------------------------|----------------------------|---------------------------|------------------------|--------------------|
| Sensitivity 0.783 | Specificity 0.55 | Precision 0.635 | Recall 0.783 | F1 0.701 |
| Accuracy 0.667 | | Kappa 0.333 | | |


```

pred_raw2 <- predict(emp_net2, testing2, decision.values=TRUE, type = "raw")

pred2 <- ROCR::prediction(pred_raw2, testing2$Attrition)
perf2 <- ROCR::performance(pred2, measure="tpr", x.measure="fpr")
plot(perf2, lwd=2, col="blue")
abline(a=0, b=1)
auc_ROCR2 <- performance(pred2, measure="auc")
auc_ROCR2 <- auc_ROCR2@y.values[[1]]
auc_ROCR2

```

```
## [1] 0.7334722
```



For neural network model result is interesting compare to SVM model. Even though the model is biased towards classifying employee as not having attrition, this result can be adjusted by moving the cutoff value. Both the full sample model and under sampled model shows similar AUC value thus choosing an appropriate cutoff value will allow the model to predict attrition of employee with desirable sensitivity and specificity.